

Deep Learning

Optimizers

Edgar F. Roman-Rangel.
`edgar.roman@itam.mx`

Digital Systems Department.
Instituto Tecnológico Autónomo de México, ITAM.

January 22nd, 2021.

Last session

- ▶ GD - SGD.
- ▶ MLP.
- ▶ Backprop.
- ▶ Multiple outputs.
- ▶ Activation functions.

Today's outline

- ▶ Code revision.
- ▶ Paper 1 discussion.
- ▶ Paper 2 discussion.
- ▶ Optimization functions.
- ▶ New code.



Outline

Contour plots

Optimization functions







Parameter space

Example of different loss functions with 2 parameters.
Seen in 3D, and from above.



Colored arrows represent different possible paths to reach the local minimum, as followed by different optimization strategies.

Several local minima



Images from Video: Contour plots - Multivariable calculus - Khan Academy
<https://www.youtube.com/watch?v=WsZj5Rb6do8>

Example for SGD



- Batch gradient descent
- Mini-batch gradient Descent
- Stochastic gradient descent

SGD Limitation

$$\omega_i = \omega_i - \nabla_{\omega_i} \mathcal{L}(y, \hat{y}).$$

- ▶ Slows down around ravines.
- ▶ Oscillates across the slopes of the ravine.
- ▶ Limited progress towards the local minimum.
- ▶ Might never escape from saddle points.

Qian, 1999. “On the momentum term in gradient descent learning algorithms”.

Outline

Contour plots

Optimization functions

Parámetros

Parámetros: w , aprendidos automáticamente durante entrenamiento

Hyperparámetros: ajustados manualmente mediante validación.

Ej: tasa aprendizaje η
optimizador
función de costo

capas y neuronas

épocas

activaciones no lineales

tamaño del lote

Notation

In this section, we will use subscripts to index time, e.g., ω_t refers to the value of parameter ω at time t .

Also, we omit the position index (i -th element) commonly seen as subscript, i.e, ω_i .

Momentum

$$\omega_t = \omega_{t-1} - v_t,$$

$$v_t = \gamma v_{t-1} + \eta \nabla_{\omega} \mathcal{L}(y, \hat{y}), \quad v_0 = 0, \quad \gamma = 0.9, \eta \approx 0.001.$$

- ▶ Accelerates SGD in the relevant direction.
- ▶ Dampens oscillations.
- ▶ Includes a fraction of the historic direction.
- ▶ Momentum accelerates for gradients pointing in the same direction, and reduces for those in changing direction.

Sutskever et al., 2013. “On the importance of initialization and momentum in deep learning”.

Nesterov Accelerated Gradient (NAG)

$$\omega_t = \omega_{t-1} - v_t,$$

$$v_t = \gamma v_{t-1} + \eta \nabla_{(\omega - \gamma v_{t-1})} \mathcal{L}(y, \hat{y}),$$

- ▶ $\nabla_{(\omega - \gamma v_{t-1})}$ approximates the next position of ω .
- ▶ Looks ahead by calculating the gradient w.r.t. future positions.
- ▶ Anticipates changes in the direction of the gradient.

Nesterov, 1983. "A method for unconstrained convex minimization problem with the rate of convergence $\mathcal{O}(1/k^2)$ ".

Adaptive Gradient (AdaGrad)

$$\omega_t = \omega_{t-1} - \eta \frac{1}{\sqrt{G_t + \epsilon}} g_t, \quad g_t = \nabla_{\omega} \mathcal{L}(y, \hat{y}),$$

$$G_t = \sum_{k=0}^t g_k^2, \quad \epsilon \approx 1e^{-8}.$$

- ▶ G_t : sum of gradients² up to time t (heavy memory loads).
- ▶ Adapts η at each time step (always decreasing).
- ▶ Works well on sparse data and large models.

Duchi et al., 2011. “Adaptive Subgradient Methods for Online Learning and Stochastic Optimization”.

Adadelta

$$\omega_t = \omega_{t-1} - \eta \frac{1}{\sqrt{\mathbb{E}[g^2]_t}} g_t, \quad g_t = \nabla_{\omega} \mathcal{L}(y, \hat{y}),$$
$$\mathbb{E}[g^2]_t = \gamma \mathbb{E}[g^2]_{t-1} + (1 - \gamma) g_t^2, \quad \gamma = 0.9.$$

- ▶ Addresses the issue of monotonically decreasing η .
- ▶ Restricts the past to a moving window.
- ▶ Recursively computes the sum of past gradients using exponential smoothing.

Zeiler, 2012. “ADADELTA: An Adaptive Learning Rate Method”.
RMSprop: a variant by Hinton (unpublished).

Adaptive Momentum (Adam)

Adadelta + Momentum.

$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$, first moment estimate (mean),
 $v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$, second moment estimate (stddv).

Correcting for bias towards zero:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}; \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}.$$

$$\omega_t = \omega_{t-1} - \eta \frac{1}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t,$$

Kingma & Ba, 2015. "Adam: a Method for Stochastic Optimization".

Other variants: AdaMax, Nadam, AMSGrad.

To know more

Ruder, 2016. "An overview of gradient descent optimization algorithms". <https://arxiv.org/abs/1609.04747>

Recette:

• Intender con Adam, SGD con momentum

Q&A

Thank you!

`edgar.roman@itam.mx`